


Clean eligibility across model-tokenizer setups



Supporting clean-baseline comparison only. No BPE or Unigram attack-generalisation experiment was executed.

Differences cannot be attributed to tokenizer family alone because the classifier architectures also differ.